

SCC.121: Fundamentals of Computer Science

Sorting, Trees and Graphs

Self-balancing Binary Search Trees

Today's Lecture

Aim:

- Introduce:
 - **Dictionary (Map)** ADT
 - Binary Search Trees (**BST**)
 - **Self-balancing** Search Trees (2-3 trees)

Searching Problem/Index System

- Set of **entries**: $\{e_1, \dots, e_n\}$
- Each entry e_i is a pair (k_i, v_i) , where:
 - k_i is a **key value** from some totally ordered domain, say integers or strings,
 - v_i is an associated **data value**.
- We assume keys are **unique**.
- We **do not use data values** to solve the searching problem, but they will be important to the applications that use our data structure.
- **Searching problem**: Given an arbitrary search key k , determine if there exists an entry matching this key value and if so, what the associated data value is.

Example of Index System

- **DNS lookup**
 - Insert domain name with specified IP address to obtain the set of (domain name, IP address) entries
 - Given domain name, find the corresponding IP address

Domain name	IP address
www.google.com	142.250.189.164
www.lancs.ac.uk	148.88.65.80
www.gov.uk	151.101.40.144
www.wikipedia.com	198.35.26.98

Examples of Index System

application	Purpose of search	key	value
dictionary	Find definition	Word	definition
Web search	Find relevant web pages	keyword	List of page names
Compiler	Find properties of variables	Variable name	Type and value
Routing table	Route internet packets	Destination	Best route
DNS	Find IP address	Domain name	IP address
Reverse DNS	Find domain name	IP address	Domain name
Genomics	Find markers	DNA string	Known positions
File system	Find file on disk	Filename	Location on disk

Dictionary (Map) ADT

- **Dictionary** (or map) is an ADT with **operations insert, delete and find**:
 - **void insert(Key k , Value v)**: Stores an entry with the key-value pair (k, v) .
(If k is already present in the dictionary, error.)
 - **void delete(Key k)**: Deletes the entry associated with key k from the dictionary.
(If it does not exist, error.)
 - **Value find(Key k)**: Determine whether there is an entry in the dictionary associated with key k . If there is, return the associated value. Otherwise, return null reference.
- You can also support **iteration** over the entries, **range queries** (enumerate or count all entries between some minimum and maximum value), **returning i th smallest value** as well as **union** and **intersection** (of dictionaries).

Dictionary (Map) ADT

- Dictionary (or map) is an ADT that supports operations **insert**, **delete** and **find**:
 - **void insert(Key k , Value v):** Stores an entry with the key-value pair (k, v) .
If k is already present in the dictionary, error.
 - **void delete(Key k):** Deletes the entry associated with key k from the dictionary.
If it does not exist, error.
 - **Value find(Key k):** Determine whether there is an entry in the dictionary associated with key k . If there is, return the associated value. Otherwise, return null reference.
- Common implementations of dictionary ADT use:
 1. Sorted arrays
 2. Search trees
 3. Hashing

Sequential allocation for Dictionary ADT

- Most naïve implementation: Store the entries in a linear array without sorting.
 - To find key value, linear search $\rightarrow O(n)$ time
 - Insertion only takes $O(1)$ time, if we still have enough space in the array.
But checking for duplicates $\rightarrow O(n)$ time.
- Otherwise, use a sorted (by key value) array.
 - To find key value, binary search $\rightarrow O(\log n)$ time.

For $n = 10^9$ entries,
 $\log_2 n \leq 30$


Sequential allocation for Dictionary ADT

- Most naïve implementation: Store the entries in a linear array without sorting.
 - To find key value, linear search $\rightarrow O(n)$ time
 - Insertion only takes $O(1)$ time, if we still have enough space in the array.
But checking for duplicates $\rightarrow O(n)$ time.
- Otherwise, use a sorted (by key value) array.
 - To find key value, binary search $\rightarrow O(\log n)$ time.
 - **But updates (insertion and deletion) take $O(n)$ time as elements in array must be moved around to make space.**


For $n = 10^9$ entries,
 $\log_2 n \leq 30$

Binary Search Tree

- Binary tree is a tree with at most two children per node: **left** child and **right** child




- Unlike the Tree ADT in the previous lecture, the following two binary trees are \neq



Binary Search Tree

- Binary search tree =
 1. Binary tree,
 2. Each node has a key-value pair,
 3. Nodes are “sorted” by their keys.
- **Sorted** = For every node, its key is:
 - Greater or equal to all keys in its left subtree,
 - Strictly smaller than all keys in its right subtree



BST ADT: Setting up

```
public class BST<Key extends Comparable<Key>,Value>{
    private Node root;

    private class Node{
        private Key key;
        private Value value;
        private Node leftChild, rightChild;


        public Node(Key key, Value value){
            this.key = key;
            this.value = value;
        }
    }
    ...
}
```

BST ADT: Implementing Find

Leverage the sortedness property:

1. If root has key, return root,
2. Else if key smaller, go left,
3. Else, go right.


Find(9):



```
public Value find(Key key){
    Node x = root;
    while (x != null){
        int cmp = key.compareTo(x.key);
        if (cmp < 0) x = x.leftChild;
        // cmp < 0 is equivalent to key < x.key
        else if (cmp > 0) x = x.rightChild;
        else return x.value; //cmp == 0
    }
    return null;
}
```

BST ADT: Time Complexity of Find

- For a BST of height h , the worst-case complexity of `find()` is $O(h)$.
- Height of BST = $\Omega(\log n)$ and $O(n)$



- BST is very efficient on balanced trees, but inefficient on degenerate trees

BST ADT: Implementing Insertion

Need to maintain sortedness:


1. If x is null, insert new node
2. If x has key, give error,
3. Else if $\text{key} < \text{x's key}$, recurse left
4. Else recurse right

```
public void insert(Key key, Value value){
    root = insert(root, key, value);
}

private Node insert(Node x, Key key, Value value){
    if (x == null) return new Node(key, value);
    int cmp = key.compareTo(x.key);
    if (cmp < 0){
        x.leftChild = insert(x.leftChild, key, value);
    }
    else if (cmp > 0){
        x.rightChild = insert(x.rightChild, key, value);
    }
    else System.out.println("Insert Exception");
        //x.value = value;
    return x;
}
```

BST ADT: Implementing Insertion


- Example: **insert(6, 'Jane')**



BST ADT: Insertion Order Matters

You can have multiple BSTs with the same key-value pairs:

1. insert(5,*)
2. insert(3,*)
3. insert(6,*)
4. insert(7,*)
5. insert(8,*)
6. insert(9,*)




Degenerate Tree

BST ADT: Insertion Order Matters

You can have multiple BSTs with the same key-value pairs:

1. insert(7,*)
2. insert(5,*)
3. insert(3,*)
4. insert(6,*)
5. insert(8,*)
6. insert(9,*)



Balanced Tree

BST ADT: Time Complexity of Insertion

Time Complexity:

- For a BST of height h , worst-case time complexity = $O(h)$
- Depends on the insertion order

How can we avoid inefficient search?

1. If **insertion order is (key) random**, then (expected) BST height is $O(\log n)$,

Average-case time complexity: $O(\log n)$

2. Otherwise, implement more complex insertion operations to “**balance**” BST height at $O(\log n)$ **for all possible insertion orders**

Self-Balancing BSTs

BST ADT: Implementing Deletion

Intuition and key steps:

1. **“Find”** the node with the value you want to delete
2. Then, there are **three possible cases**:
 - a) Node is **leaf**



BST ADT: Implementing Deletion

Intuition and key steps:

1. **“Find”** the node with the value you want to delete
2. Then, there are **three possible cases**:
 - a) Node is **leaf**
 - b) Node has a **single child**



BST ADT: Implementing Deletion

Intuition and key steps:

1. **“Find”** the node with the value you want to delete
2. Then, there are **three possible cases**:
 - a) Node is **leaf**
 - b) Node has a **single child**
 - c) Node has **both children**

- Find “successor” of 3,
- Use it to replace
- Ensures:
 - keys in left subtree remain smaller,
 - keys in right subtree greater.



Binary Search Tree: Code for delete()

Successor of a node with value = node in tree with value just above that

```
private Node successor(Node x){  
    if (x == null || x.rightChild == null) return null;  
    else{  
        x = x.rightChild;  
        while(x.leftChild != null) x = x.leftChild;  
        return x;  
    }  
}
```

Binary Search Tree: Code for delete()

```
public void deleteMin(){
    root = deleteMin(root);
}

private Node deleteMin(Node x){
    if (x.leftChild == null) return x.rightChild;
    else {
        x.leftChild = deleteMin(x.leftChild);
        return x;
    }
}


public void delete(Key key){
    root = delete(root, key);
}
```


Binary Search Tree: Code for delete()

```
private Node delete(Node x, Key key){
    if (x == null) return null;
    int cmp = key.compareTo(x.key); // cmp = sign(key - x.key) for integers
    if (cmp < 0) x.leftChild = delete(x.leftChild, key);
    else if (cmp > 0) x.rightChild = delete(x.rightChild, key);
    else{ // if we need to delete node x
        if (x.rightChild == null) return x.leftChild;
        else if (x.leftChild == null) return x.rightChild;
        else{
            Node t = x;
            x = successor(t);
            x.rightChild = deleteMin(t.rightChild);
            x.leftChild = t.leftChild;
        }
    }
    return x;
}
```

Beyond Simple Binary Search Trees

- All operations on a BST of height h have a worst-case time complexity of $O(h)$.
- But that will also hold for non-binary search trees.
 - Where you have $i \leq \max$ keys per tree node and $i + 1$ children.
 - The first two subtrees contain, respectively, nodes whose keys are smaller and greater than the first key.
 - The second and third subtree whose keys are smaller and greater than the second key.




- To make these operations more efficient, use **self-balancing search trees**

Self-Balancing Search Trees

- There are many options, with diverse tradeoffs
 - **2-3 trees**: tertiary self-balancing search trees
 - AVL trees: self-balancing BST
 - Red-black trees are a (BST) generalization of 2-3 trees
 - B-trees: generalization of 2-3 trees to more keys per node
- Some are **self-balancing but not height-balanced**:
 - Splay Trees
 - Treaps
- For example, you may want your search tree to self-balance for **common letter (key) patterns** rather than for any letter (key) sequence.

2-3 Trees


- 2-nodes: 1 key, 2 children
- 3-nodes: 2 keys, 3 children



- (Natural variant of) Inorder traversal gives keys in ascending order
- **Perfectly balanced tree:**
Every path from the root to a leaf has the same length
- How do we maintain this property despite insertions and deletions?

2-3 Trees: Find



- Works (almost) just as in BSTs,
- Example: **find(8)**





2-3 Trees: Insertion

Slightly more complex than in BSTs. **Three cases:**

1. insert(5)

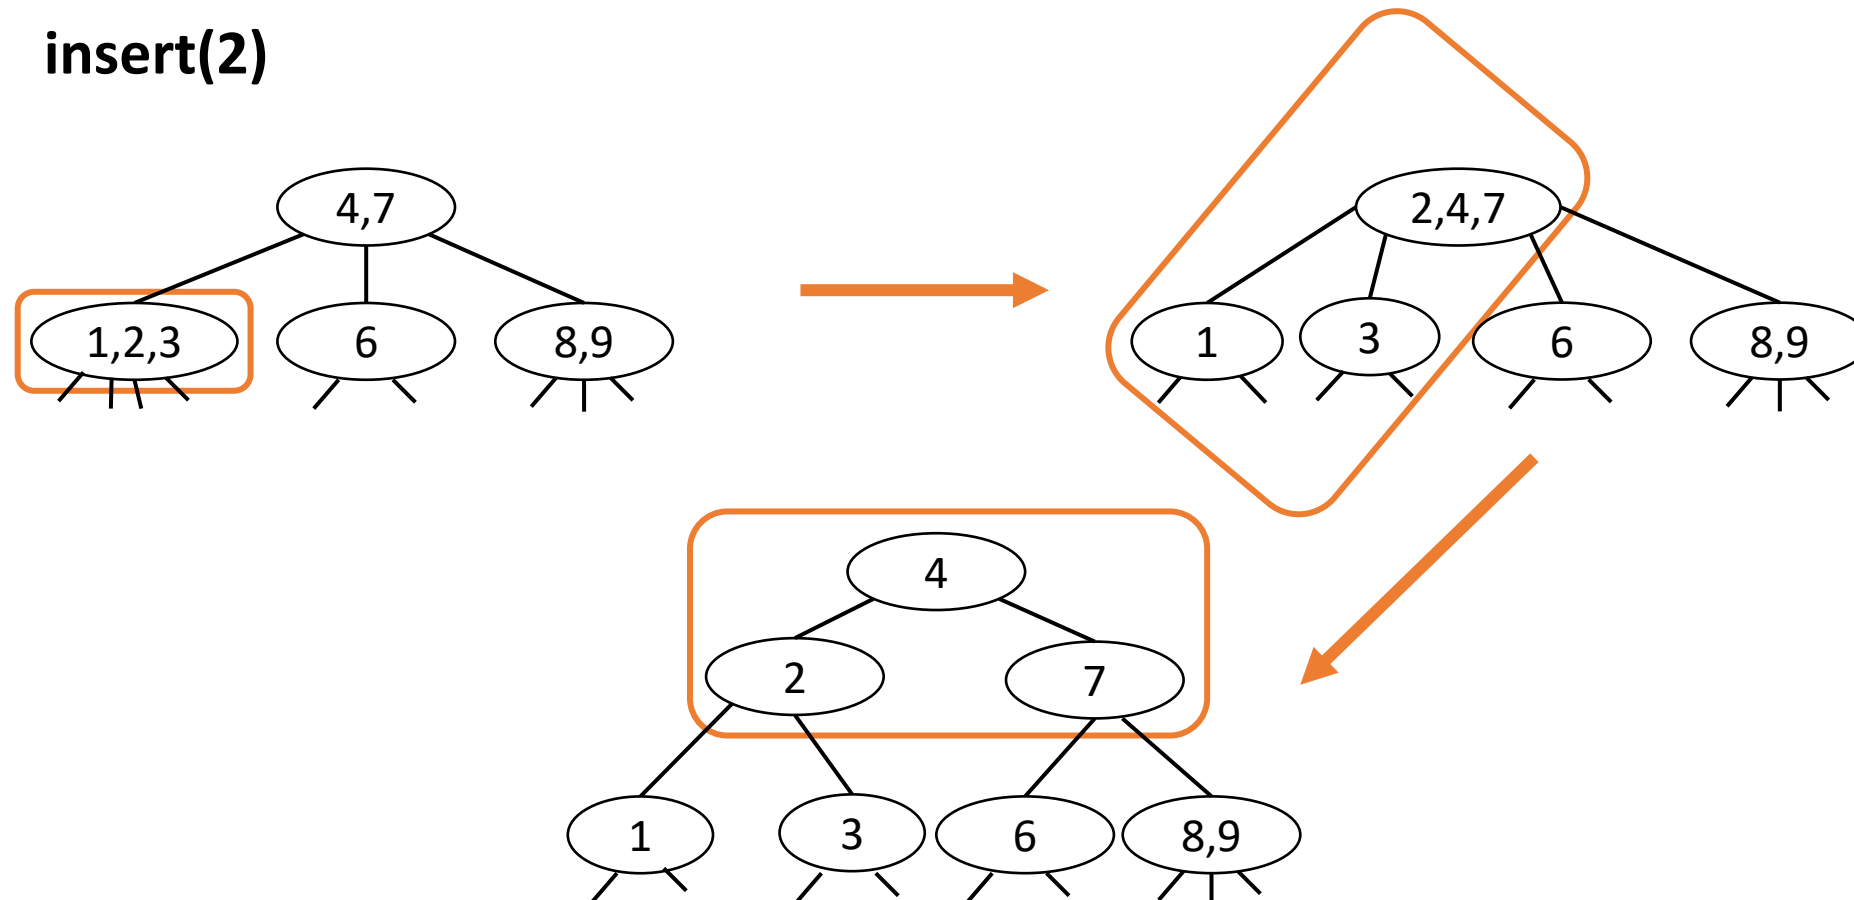


2. insert(2)




2-3 Trees: Insertion

3. insert(2)



2-3 Trees: Insertion

- Why does insertion maintain perfect balance?
 - Either a 2-node is transformed into a 3-node,




2-3 Trees: Insertion

- Why does insertion maintain perfect balance?
 - Either a 2-node is transformed into a 3-node,
 - Or a 3-node is split into three 2-nodes,




2-3 Trees: Insertion

- Why does insertion maintain perfect balance?
 - Either a 2-node is transformed into a 3-node,
 - Or a 3-node is split into three 2-nodes (root),
 - Or a (non-root) 3-node with a 2-node parent transforms into a temporary 4-node,




2-3 Trees: Insertion

- Why does insertion maintain perfect balance?
 - Either a 2-node is transformed into a 3-node,
 - Or a 3-node is split into three 2-nodes (root),
 - Or a (non-root) 3-node with a 2-node parent transforms into a temporary 4-node,
 - Or a (non-root) 3-node with a 3-node parent transforms into a temporary 4-node,



2-3 Trees: Cons

- However, 2-3 trees are quite inconvenient to implement.
 1. Different types of nodes (2-nodes and 3-nodes)
 2. We need to do multiple compares at each node,
 3. We need to move back up the tree to split 4-nodes
 4. Large number of cases for that splitting.
- Red-black trees are BSTs that build upon 2-3 trees, representing 3-node as a binary tree node with a **red left link**



Summary

Today's lecture:

Introduced:

- Dictionary (Map) ADT,
 - Binary Search Trees
 - And Self-balancing Search Trees (2-3 trees)
-
- **Next Lecture:** Graphs and Shortest Paths
 - **Any questions?**