

SCC.131: Digital Systems Debugging

Ioannis Chatzigeorgiou (i.chatzigeorgiou@lancaster.ac.uk)

Partially based on material produced by Charalampos Rotsos

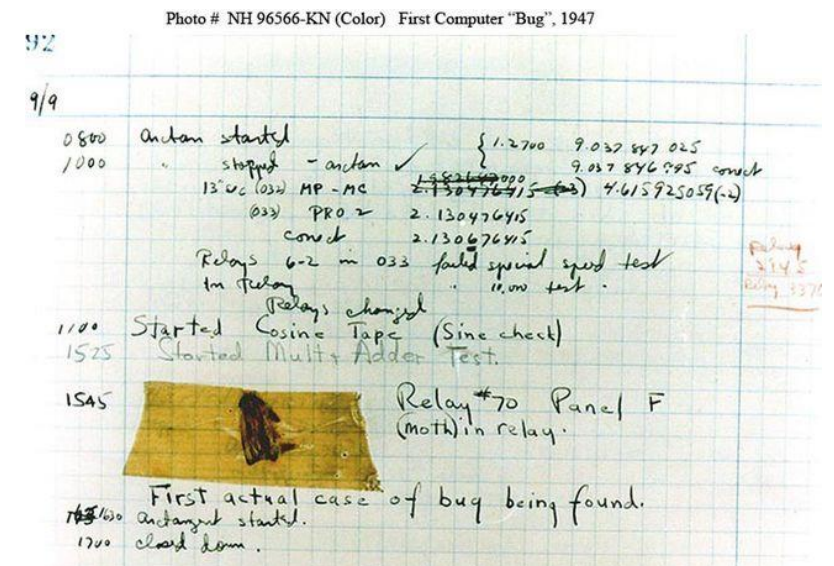
Summary of the last lecture

The following points were covered in the last lecture:

- Initial processing: merge continued lines, break them, remove comments.
- Tokenization: each line is broken down into 'tokens'; the preprocessor looks at tokens that contain directives and macros.
- The flexibility that preprocessing offers in:
 - Replacing object-like and function-like identifiers (names of macros) with their definitions (bodies of macros).
 - Including header files and controlling compilation (using conditionals and computed include directives).
 - Using predefined macros to diagnose problems.

The first 'bug'

- Debugging is a methodical process of finding and reducing the number of bugs, or defects, in a computer program.
- In September 1947, a team led by Dr Grace Hopper at Harvard University traced an error in the Harvard Mark II computer to a moth trapped in a relay, coining the term 'bug'.
- This bug was carefully removed and taped to the logbook. Stemming from the first bug, today we call errors or glitches in a program a 'bug'.



Types of bugs

- Bugs that occur during **compilation time**.
 - Your program code is syntactically incorrect.
 - Your program violates common programming conventions, for example a variable is used before / without initialization.
 - Issues may appear as warnings.
 - Static analysis of your code detects that the program is invalid.
- Bugs that occur during **run time**.
 - Your program has a logical error.
 - It will still work but not as expected.

Spot the three errors

```
#include "MicroBit,h"

MicroBit uBit;

/* Print "hello, world" to stdout and
return 0.

int main(void)
{
    uBit.init();
    uBit.display.scroll("HELLO WORLD\n")
    return 0;
}
```

Spot the three errors

```
#include "MicroBit.h"

MicroBit uBit;


/* Print "hello, world" to stdout and
return 0. */

int main(void)
{
    uBit.init();
    uBit.display.scroll("HELLO WORLD\n");
    return 0;
}
```


Debugging strategies

- **Divide and conquer** to debug a program:
 - Gradually remove/add code to create the smallest source file that contains the bug.

- **Approach 1: Remove code**
 - Start with your code
 - Slowly remove code until program work well.
 - Examine last removed lines.



- **Approach 2: Add code**
 - Start with the smallest working program.
 - Add functionality, until program breaks.
 - Examine the last added lines



Further debugging using logging (`printf`)

Another approach is to insert `printf` in different places of the code to follow its flow. This approach is often helpful, but:

- It generally takes an awful lot of `printf` statements.
- Inclusion of `printf` could change the behaviour of the code (timing, stack, ...).
 - programs with `printf` may work, and then ...
 - ... the code fails when `printf` is removed (this is known as a 'heisenbug').
- This approach cannot examine program flow details, i.e., instruction-by-instruction inspection.

Example of issue hidden by printf

- This program could compile and run but contains a nasty bug! **Can you spot it?**

```
#include <stdio.h>
```

```
void test(){  
    int k;  
    *(&k+288)=3;  
}
```

```
int main() {  
    int i = 5;  
    char a[10]="TEST1";  
    printf("%s\n",a);  
    test();  
    i=i+i;  
    printf("%d\n",i);  
}
```

Example of issue hidden by printf

- This program could compile and run but contains a nasty bug! **Can you spot it?**
- When all printf statements are commented out (or removed), the stack will be affected, i.e., memory addresses and memory allocation will change, and a bus error could be triggered.
- Also depends on the C compiler (with/without optimisation).

```
#include <stdio.h>
```

```
void test(){  
    int k;  
    *(&k+288)=3;  
}
```

```
int main() {  
    int i = 5;  
    char a[10]="TEST1";  
    printf("%s\n",a);  
    test();  
    i=i+i;  
    printf("%d\n",i);  
}
```

Micro:bit and printing on the screen

- You can use `printf` and `scanf` operations to print and read data using a **serial/UART** interface.
 - A communication interface between two computers.
 - Transmits information sequentially one bit at a time.
 - You need a serial client to read and write data (i.e., screen).



Micro:bit: Sending data through the serial (1/2)

- For **Linux** systems:
 - Install the program “screen” if it is not already installed.
 - Upload your code to your micro:bit device and open a terminal window.
 - Type `ls /dev/ttyACM*` to find out the device node that micro:bit has been assigned to. We are going to assume that the device node is `/dev/ttyACM0`.
 - Type `screen /dev/ttyACM0 115200` to display the serial output of micro:bit on your screen. The value 115200 is the default baud rate (symbols/sec) of micro:bit.
 - To exit “screen”, press `Ctrl-A` and `Ctrl-D`. To return to “screen” type `screen -r`
- For **OSX** systems:
 - Same as Linux but use `/dev/cu.usbmodem*` instead of `/dev/ttyACM*`.

Micro:bit: Sending data through the serial (2/2)

- For **Windows** systems that use the Windows Subsystem for Linux (WSL):
 - Open a terminal window as an **administrator**.
 - Install usbipd:

```
winget install --interactive --exact dorssel.usbipd-win
```
 - Type `usbipd list` to find out the bus IDs of currently connected USB devices.
Let us assume that micro:bit uses bus 2-7.
 - Bind and then attach to the device through WSL:

```
usbipd bind -b 2-7  
usbipd attach -b 2-7 -w -a
```
 - Open a **WSL window** as an **administrator** and follow the guidelines for Linux.

Micro:bit: Debugging using logging (printf)

- The `MicroBitSerial` class (`uBit.serial.*`) allows you to call `printf` and `scanf` functions, as you can see in this example.
- Upload this code to your micro:bit **before** you use the `screen` command, and you will see only the message "HELLO WORLD!" scrolling across the display.
- When you run `screen`, the serial output will be displayed on the screen of your PC and the message "MICROBIT SAYS HELLO!" will be printed.

```
#include "MicroBit.h"

MicroBit uBit;

int main()
{
    uBit.init();
    while(1)
    {
        uBit.display.scroll("HELLO WORLD!", 100);
        uBit.serial.printf("MICROBIT SAYS HELLO!\n");
    }
}
```

Using a debugger

- A debugger is the (less invasive) alternative to `printf`.
- Allows you to:
 - Step through a program (execute one instruction at a time).
 - Set breakpoints (stop at checkpoints).
 - Investigate machine state (memory, registers).
 - Investigate crashes.
- It does **not**:
 - Find problems for you (but it makes this job easier).
 - Fix a problem (you have to do that...).

Debuggers: One size fits all?

- Debuggers are generally **language dependent**.
 - Some debuggers can handle many different languages.
- Some debuggers/debugging requires **hardware** support.
 - In-System programming of logic devices, e.g., FPGAs (Field Programmable Gate Arrays) using Verilog or other Hardware Description Languages (HDLs).
 - e.g., JTAG (Joint Test Action Group) to access debug interfaces.
 - Hardware support for code/data breakpoints (page fault).
- Debuggers may provide **different interfaces**.
 - Command line.
 - Graphical user interface (GUI).

The GNU DeBugger (GDB)

- In SCC.131, we use GDB.
 - Open-source debugger developed by the GNU project that also created the GNU Compiler Collection (GCC or gcc).
 - Designed for the C language.
 - Command line interface but can be used with Integrated Development Environments (IDEs).
- Aims
 - Debugging C programs.
 - Understanding of system architecture.
 - Connection between hardware, assembler, C, applications.

Learning GDB

- GDB resources:
 - [GDB manual](#)
 - [Cheatsheet](#)
- To learn GDB, re-run C code of this module - and other SCC modules - given to you in lectures and lab sessions.
- Built-in GDB help:
 - `gdb --help`
 - `gdb help command`

Debug symbols

- During compilation, you need to ask your compiler to generate and embed **debug symbols**.
 - Records associating code and variables with source code.
 - The flag `-g`, tells gcc to generate debug symbols.
 - To compile and then debug the program on this slide, type:

```
$ gcc -g SCC131_W10_Debug.cpp -o SCC131_W10_Debug
```

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Desired output

ASCII code:

Dec	Hex	Char
65	41	A
...
80	50	P

Dec	Hex	Char
97	61	a
...
112	70	p

Equivalent to:

`s[i]='A'+i;`
`t[i]='a'+i;`

Desired program output: `>ABCDEFGHJKLMNOP<`
`>abcdefghijklmnop<`

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Actual output

- Run GDB:

```
$ gdb SCC131_W10_Debug
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04.2) 12.1
Reading symbols from SCC131_W10_Debug...
(gdb) run
Starting program: SCC131_W10_Debug
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/
libthread_db.so.1".
>ABCDEFGHJKLMNOP<
><
[Inferior 1 (process 7412) exited normally]
(gdb) █
```

What happened to the output of the second printf?

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Pause execution

- Add a breakpoint and run again:

```
(gdb) break 13
Breakpoint 1 at 0x555555551d4: file SCC131_W10_Debug.c
pp, line 13.
(gdb) run
Starting program: SCC131_W10_Debug
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/
libthread_db.so.1".
```

```
Breakpoint 1, main () at SCC131_W10_Debug.cpp:13
```

```
13         printf(">%s<\n", s);
```

```
(gdb) █
```

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Check the memory (1/3)

- List 40 bytes of memory from the address of s:

```
(gdb) x/40xb &s
```

```
0x7fffffffdf10: 0x41 0x42 0x43 0x44 0x45 0x46 0x47 0x48
0x7fffffffdf18: 0x49 0x4a 0x4b 0x4c 0x4d 0x4e 0x4f 0x50
0x7fffffffdf20: 0x00 0x62 0x63 0x64 0x65 0x66 0x67 0x68
0x7fffffffdf28: 0x69 0x6a 0x6b 0x6c 0x6d 0x6e 0x6f 0x70
0x7fffffffdf30: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
```

```
(gdb) █
```

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Check the memory (2/3)

- What does this mean?

Memory dump showing addresses and values:

Address	0x41	0x42	0x43	0x44	0x45	0x46	0x47	0x48
0x7fffffffdf10:	0x41	0x42	0x43	0x44	0x45	0x46	0x47	0x48
0x7fffffffdf18:	0x49	0x4a	0x4b	0x4c	0x4d	0x4e	0x4f	0x50
0x7fffffffdf20:	0x00	0x62	0x63	0x64	0x65	0x66	0x67	0x68
0x7fffffffdf28:	0x69	0x6a	0x6b	0x6c	0x6d	0x6e	0x6f	0x70
0x7fffffffdf30:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00

Annotations:

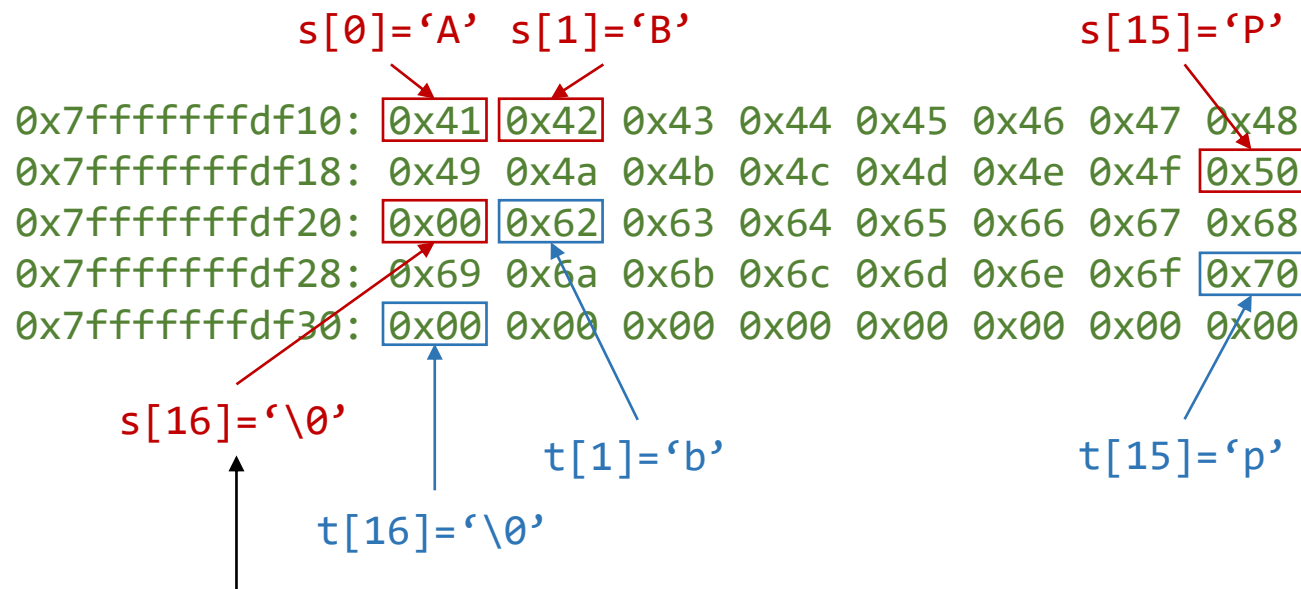
- $s[0] = 'A'$ (points to 0x41)
- $s[1] = 'B'$ (points to 0x42)
- $s[15] = 'P'$ (points to 0x50)
- $s[16] = '\0'$ (points to 0x00 at address 0x7fffffffdf20)

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```


GDB – Example – Check the memory (3/3)

- What does this mean?



The starting address of `t` is `t[0]`, which has been overwritten by `s[16]` and now contains `'\0'`.

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[16], t[16];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Example – Fix the bug

- Memory allocation after the bug is fixed:

Memory dump showing the state of memory after the bug is fixed:

0x7fffffffdef0:	0x41	0x42	0x43	0x44	0x45	0x46	0x47	0x48
0x7fffffffdef8:	0x49	0x4a	0x4b	0x4c	0x4d	0x4e	0x4f	0x50
0x7fffffffdf00:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffffdf08:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00
0x7fffffffdf10:	0x61	0x62	0x63	0x64	0x65	0x66	0x67	0x68
0x7fffffffdf18:	0x69	0x6a	0x6b	0x6c	0x6d	0x6e	0x6f	0x70
0x7fffffffdf20:	0x00	0x00	0x00	0x00	0x00	0x00	0x00	0x00

Annotations:

- $s[0] = 'A'$ (points to 0x41)
- $s[15] = 'P'$ (points to 0x50)
- $s[16] = '\backslash 0'$ (points to 0x00 at 0x7fffffffdf00)
- $t[0] = 'a'$ (points to 0x61)
- $t[15] = 'p'$ (points to 0x70)
- $t[16] = '\backslash 0'$ (points to 0x00 at 0x7fffffffdf20)

SCC131_W10_Debug.cpp

```
1  #include <stdio.h>
2
3  int main()
4  {
5      char s[17], t[17];
6      int i=0;
7      for(i=0; i<16; i++) {
8          s[i]=65+i;
9          t[i]=97+i;
10     }
11     s[i] = '\0';
12     t[i] = '\0';
13     printf(">%s<\n", s);
14     printf(">%s<\n", t);
15 }
```

GDB – Another example – Byte order

- Each **byte** of memory is given a **unique address**.
- In the previous example, a single character occupies exactly 1 byte, so **each** character of a string is allocated a unique address.
- What happens when a variable occupies more than 1 byte, as in the example on this slide?
- Notice that 65000 is equivalent to **0xfde8** (2 bytes).
- Repeat the same process (i.e., add debug symbols, run gdb, insert a break in line 9, run the program) and type:

(gdb) x/2xb &s

0x7fffffffdf38: 0xe8 0xfd

(gdb) █



↑ Why is the **least** significant byte displayed first?

SCC131_W10_Store2Bytes.cpp


```
1  #include <stdio.h>
2
3  int main()
4  {
5      int s;
6
7      s = 65000;
8
9      printf("s=%d\n", s);
10 }
```

Big Endian vs. Little Endian

- Jonathan Swift's **Gulliver's Travels**
 - **Big Endians** broke their eggs on the big end of the egg
 - **Little Endians** broke their eggs on the little end of the egg



BIG ENDIAN - The way people always broke their eggs in the Lilliput land




LITTLE ENDIAN - The way the king then ordered the people to break their eggs

BIG ENDIANS fled Lilliput and gained favour in Blefuscu.

Those who stayed in Lilliput became **LITTLE ENDIANS**.

Big Endian vs. Little Endian

LITTLE ENDIAN - The Lilliputians break their eggs at the smaller end




Little-endian

32-bit integer


0A0B0C0D

The Least Significant Byte (LSB) is stored first

Byte-ordered Memory



BIG ENDIAN - The Blefuscudians break their eggs at the big end.



Big-endian

32-bit integer

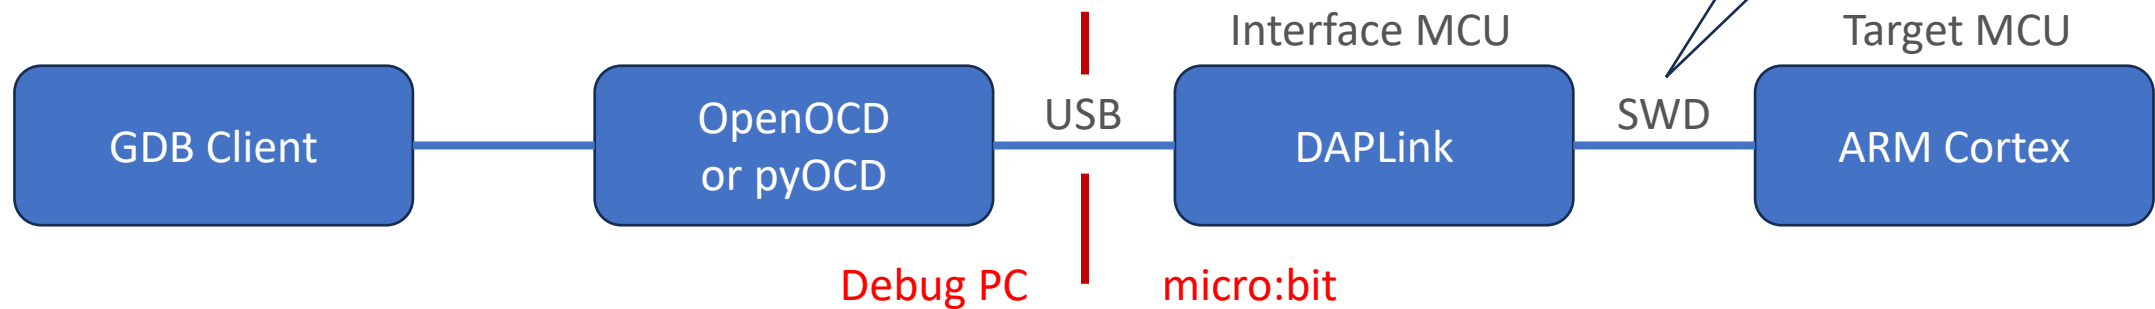
0A0B0C0D

The Most Significant Byte (MSB) is stored first

Debugging micro:bit

- We need to use **DAPLink** to debug a running micro:bit.
- An **On-Chip Debugger** (OCD) allows remote debugging.

See
Week 7 –
Lecture 2



- Integration with **VS Code** provides a powerful debugging environment.

Breakpoints

Debug Controls

Debug View

Debug Console

Debug Prompt

```
source > main.cpp > main()
You, 1 minute ago | 3 authors (You and others)
1 #include "MicroBit.h"
2 MicroBit uBit;
3
4 int main () {
5     char s[16], t[16];
6     int i;
7
8     uBit.init();
9
10    for (i=0; i < 16; i++) {
11        s[i] = 'a' + i;
12        t[i] = 'A' + i;
13    }
14    s[i] = '\0';
15    t[i] = '\0';
16    uBit.serial.printf(">%s<\n", s);
17    uBit.serial.printf(">%s<\n", t);
18 }
19
```



RUN AND DEBUG micro:bit PyOCD Co

VARIABLES

- Local
 - s: [16]
 - t: [16]
 - i: <optimized out>
- Global
 - Static: ./source/main.cpp
 - Registers

PROBLEMS OUTPUT TERMINAL PORTS GITLENS MEMORY XRTOS DEBUG CONSOLE

Filter (e.g. text, lexclu...)

```
warning: Source file is more recent than executable.
14     s[i] = '\0';
→ list
9
10     for (i=0; i < 16; i++) {
11         s[i] = 'a' + i;
12         t[i] = 'A' + i;
13     }
14     s[i] = '\0';
15     t[i] = '\0';
16     uBit.serial.printf(">%s<\n", s);
17     uBit.serial.printf(">%s<\n", t);
18 }
{"output":"","token":30,"outOfBandRecord":{},"resultRecords":{"resultClass":"done","results":[]}}
→ list
Line number 19 out of range; /Users/cr409/Library/CloudStorage/Dropbox/code/microbit-v2/source/main.cpp has 18 lines.
Line number 19 out of range; /Users/cr409/Library/CloudStorage/Dropbox/code/microbit-v2/source/main.cpp has 18 lines. (from interpreter-exec console "list")
> |
```

WATCH

CALL STACK

BREAKPOINTS

- ☒ main.cpp source 10
- ☐ test.s source 17

CORTEX LIVE WATCH

XPERIPHERALS

PERIPHERALS

REGISTERS

MEMORY

DISASSEMBLY

Summary

Today we learnt about:

- Types of bugs and debugging strategies.
- Debugging using logging (`printf`) and reading the serial port to display messages sent by micro:bit.
- Potential issues that could arise when `printf` is removed.
- Debugging using GDB.
- Big-endian and little-endian systems.
- On-Chip Debugging (OCD).

Resources

- Harvard Mark II: https://en.wikipedia.org/wiki/Harvard_Mark_II
- Sep 9, 1947 CE: World's First Computer Bug:
<https://education.nationalgeographic.org/resource/worlds-first-computer-bug/>
- Heisenbug: <https://en.wikipedia.org/wiki/Heisenbug>
- GDB manual: <https://sourceware.org/gdb/current/onlinedocs/gdb.pdf>