

# SCC.131: Digital Systems

## Memory layout (focus on x86-64)

---

Ioannis Chatzigeorgiou  
([i.chatzigeorgiou@lancaster.ac.uk](mailto:i.chatzigeorgiou@lancaster.ac.uk))

# Summary of the last lecture

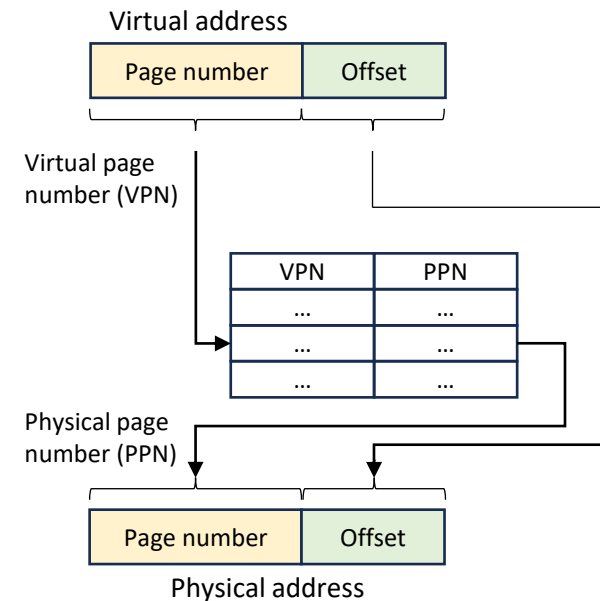
---

In the last lecture, we discussed about:

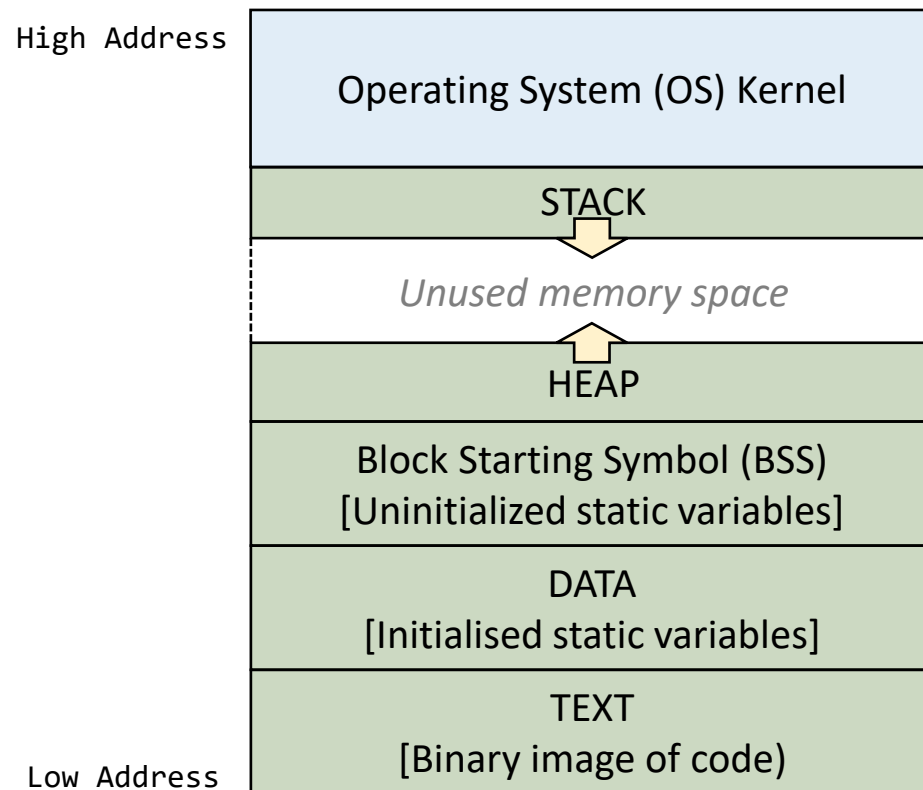
- The capabilities of micro:bit for wireless communication.
- The characteristics of the proprietary radio mode at 2.4 GHz.
- The concept of a datagram and key instructions for transmitting and receiving datagrams using variables of type PacketBuffer.
- Setting up a listener for the radio component and calling an event handler when a datagram is received.
- The creation of groups and the extraction of the received signal strength.
- Steps on how to build a simple one-way wireless communication system.

# Virtual memory

- In principle, a 64-bit microprocessor can address  $2^{64}$  bytes of byte-addressable memory.
- Read/write operations use **virtual** memory addresses.
- The **memory management unit** (MMU) in the microprocessor uses a lookup table, called the **translation lookaside buffer** (TLB), to translate virtual to physical addresses.
- The operating system dynamically allocates **pages** of memory (typically 4 KB per page for x86 architectures) to processes and creates entries in the TLB. Your code remains agnostic to coordination with other process when accessing memory.

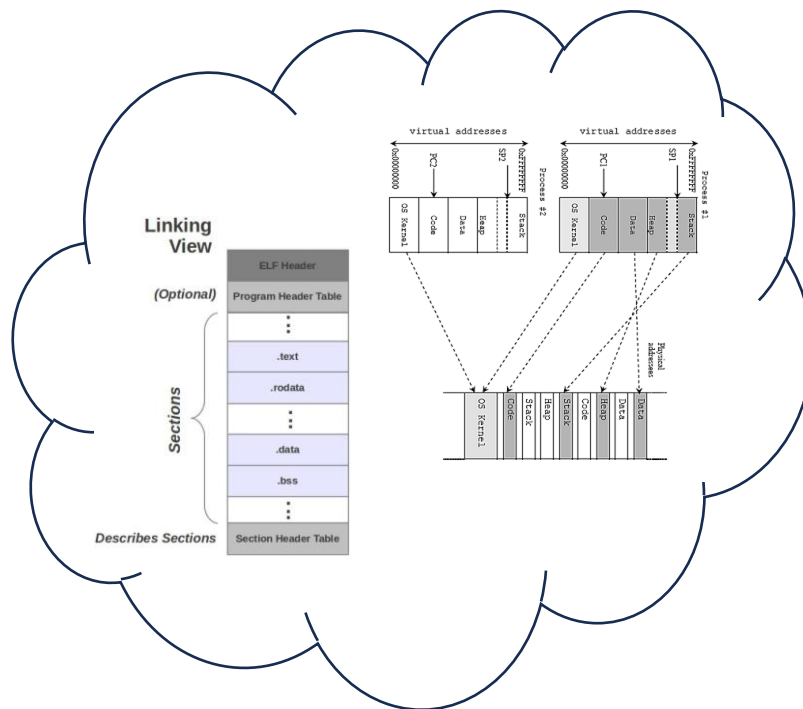


# Virtual memory layout




- **OS Kernel:** Memory reserved by the OS to monitor and control mapping between physical and virtual addresses.
- **STACK:** Data needed by function calls, including arguments and local variables. The set of values pushed for one function call is referred to as a *stack frame*.
- **HEAP:** Dynamic memory allocation for variables whose size is known at runtime and cannot be statically determined by the compiler before program execution (e.g., `malloc`, `new`, `free`, `delete`).
- **BSS** and **DATA:** Global variables that are initialized to zero or do not have explicit initialization (BSS), or global variables that have been initialized (DATA).
- **TEXT:** The binary executable instructions of a program.

# Did I hear about virtual memory before?



Week 9, Lecture 1, Slides 6 and 18  
(Compiler, assembler, linker and loader)



```
#include <stdio.h>

void test(){
    int k;
    *(&k+288)=3;
}

int main() {
    int i = 5;
    char a[10]="TEST1";
    printf("%s\n",a);
    test();
    i=i+i;
    printf("%d\n",i);
}
```

Week 10, Lecture 1, Slide 14  
(Debugging)

# Virtual memory allocation – Examples (1/2)

- Compile C program to generate object file:

```
gcc -c SCC131_example.c
```

- List the size (in bytes) of each section of the object file:

```
size SCC131_example.o
```

text	data	bss	dec	hex	filename
130	0	0	130	82	SCC131_example.o

130      82  
Total size (in decimal and hexadecimal representation)

SCC131\_example.c

```
#include <stdio.h>
int main(void) {
    printf("hello world\n");
}
```

# Virtual memory allocation – Examples (2/2)

```
#include <stdio.h>
int i=5;
int main(void) {
    printf("hello world\n");
}
```

text	data	bss	dec	hex	filename
130	4	0	134	86	SCC131_example.o

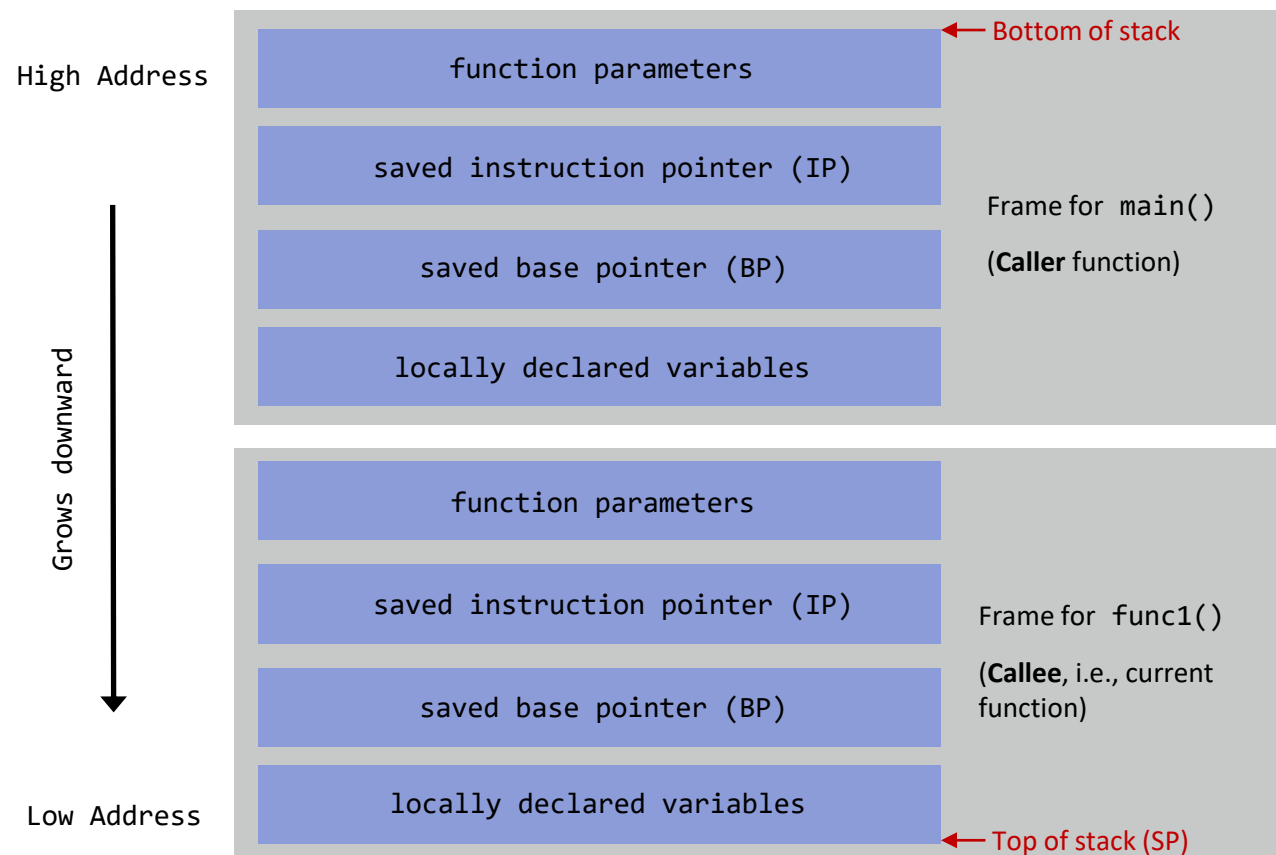
```
#include <stdio.h>
int i;
int main(void) {
    printf("hello world\n");
}
```

text	data	bss	dec	hex	filename
130	0	4	134	86	SCC131_example.o

```
#include <stdio.h>
int main(void) {
    int i=5;
    printf("hello world\n");
}
```

text	data	bss	dec	hex	filename
141	0	0	141	8d	SCC131_example.o

# Stack layout



The **stack pointer** (SP) points to the lowest address of the stack for the current frame (top of stack).

The **frame pointer** or **base pointer** (FP or BP) points to a reference address of the previous frame.

The **saved instruction pointer** (IP) points to the address that the function will return to.

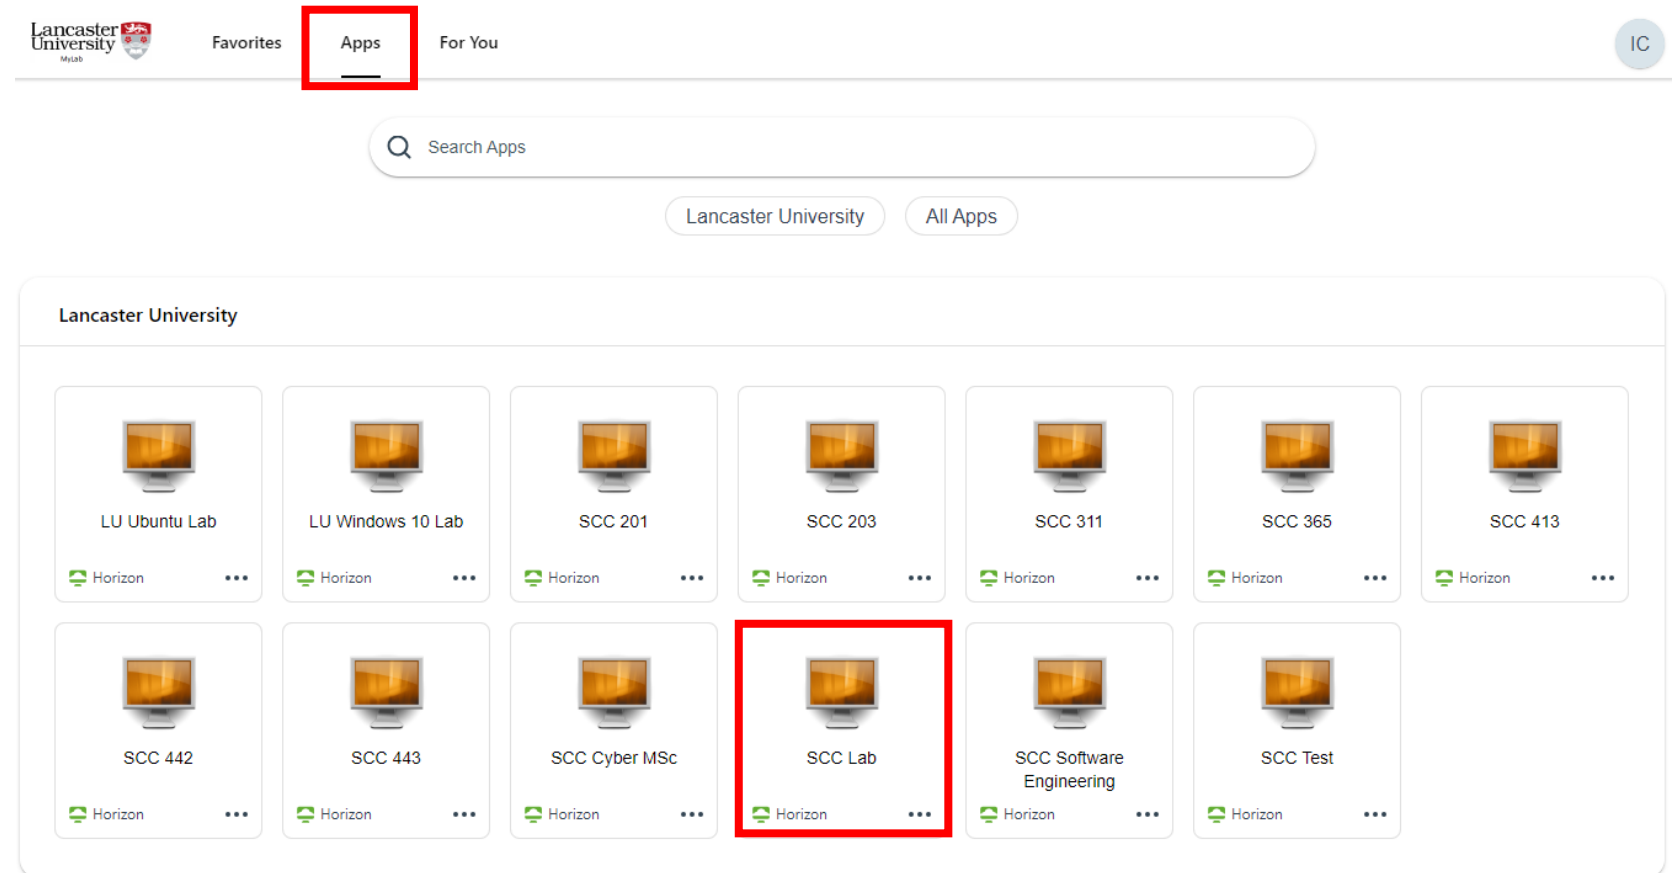
Parameter values are pushed in *reverse order*, so that the first parameter will be the last to be pushed to the stack. Therefore, the first parameter will be at a *fixed offset* from the BP regardless of the number of parameters.



# Step-by-step demonstration

Connect to [MyLab](#)  
(for guidelines,  
check [ASK](#)).

Select the 'Apps'  
tab and click 'SCC  
Lab' to access the  
appropriate virtual  
machine.



The screenshot displays the Lancaster University MyLab interface. At the top, there is a navigation bar with the Lancaster University logo, a 'Favorites' tab, an 'Apps' tab (highlighted with a red box), and a 'For You' tab. To the right of the navigation bar is a user profile icon labeled 'IC'. Below the navigation bar is a search bar labeled 'Search Apps'. Underneath the search bar are two buttons: 'Lancaster University' and 'All Apps'. The main content area is titled 'Lancaster University' and displays a grid of application tiles. Each tile features a monitor icon, a title, and a 'Horizon' logo with a three-dot menu. The tiles are arranged in two rows. The first row contains: 'LU Ubuntu Lab', 'LU Windows 10 Lab', 'SCC 201', 'SCC 203', 'SCC 311', 'SCC 365', and 'SCC 413'. The second row contains: 'SCC 442', 'SCC 443', 'SCC Cyber MSc', 'SCC Lab' (highlighted with a red box), 'SCC Software Engineering', and 'SCC Test'.

Application	Horizon
LU Ubuntu Lab	...
LU Windows 10 Lab	...
SCC 201	...
SCC 203	...
SCC 311	...
SCC 365	...
SCC 413	...
SCC 442	...
SCC 443	...
SCC Cyber MSc	...
SCC Lab	...
SCC Software Engineering	...
SCC Test	...

# Step 1: Determine the CPU architecture

---

To display information about the CPU architecture:

`lscpu`

The output will look like:

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Address sizes:      45 bits physical, 48 bits virtual
Byte Order:         Little Endian
```

# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture:

CPU op-mode(s):

Address sizes:

Byte Order:

`x86_64`

32-bit, `64-bit`

45 bits physical, 48 bits virtual

Little Endian

64-bit architecture typically used  
by Intel and AMD processors.

$2^{64}$  bytes of byte-addressable  
memory (in theory).

# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Address sizes:      45 bits physical, 48 bits virtual
Byte Order:        Little Endian
```

64-bit addresses have been  
limited to 48-bit addresses.


# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Address sizes:      45 bits physical, 48 bits virtual
Byte Order:        Little Endian
```



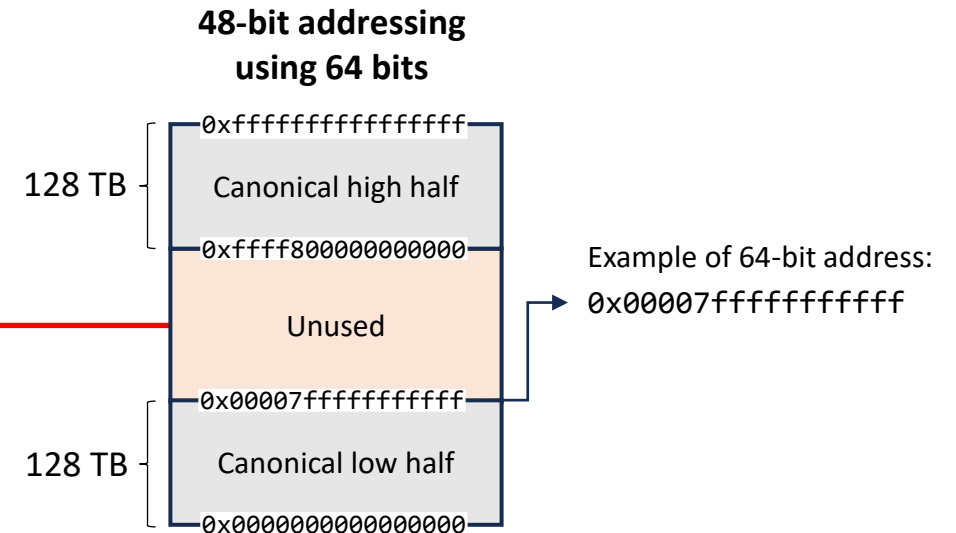
# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian




# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



Prefix used to denote  
hexadecimal (base-16)  
numbers.

0x00007fffffffffffffff

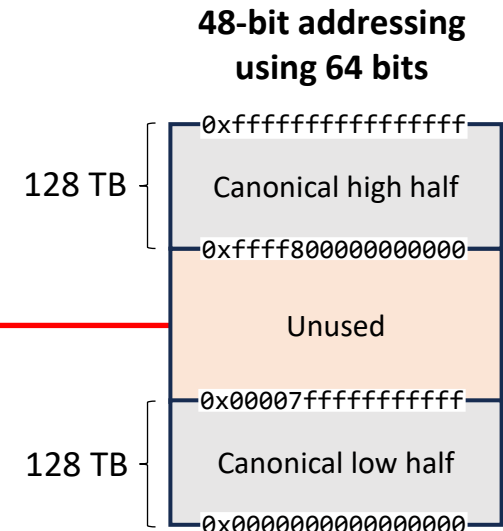
# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



Prefix used to denote  
hexadecimal (base-16)  
numbers.

Hexadecimal digit  
(16 values, i.e., 4 bits).

0x00007fffffffffffffff




# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



Prefix used to denote  
hexadecimal (base-16)  
numbers.

Hexadecimal digit  
(16 values, i.e., 4 bits).

0x00007fffffffffffffff  
4 bits × 16 digits = 64 bits


# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



Prefix used to denote hexadecimal (base-16) numbers.

Hexadecimal digit  
(16 values, i.e., 4 bits).

0x00007fffffffffffffff

4 bits × 16 digits = 64 bits

Truncate  
address

0x00007fffffffffffffff

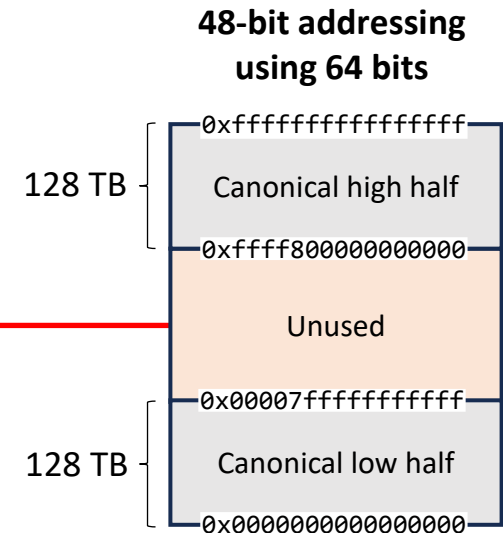
# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



Prefix used to denote hexadecimal (base-16) numbers.

Hexadecimal digit  
(16 values, i.e., 4 bits).

0x00007fffffffff

4 bits × 16 digits = 64 bits

Truncate  
address

0x00007fffffffff

4 bits × 12 digits = 48 bits

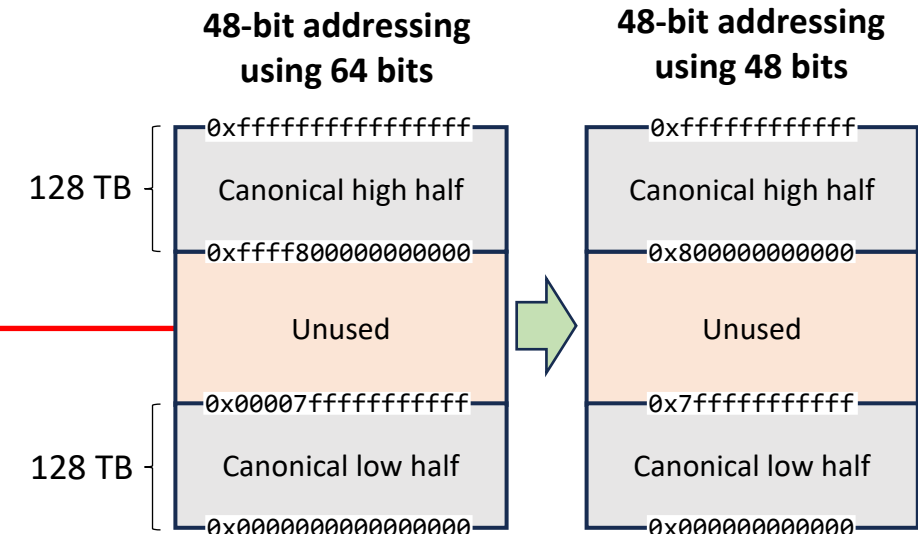
# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: Little Endian



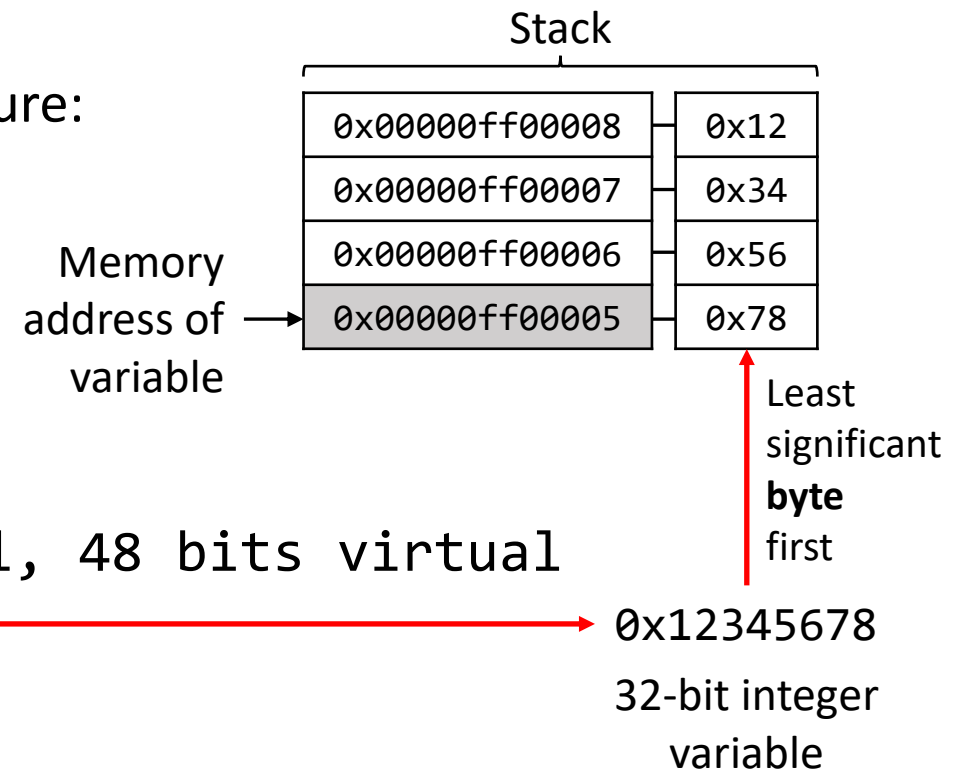
# Step 1: Determine the CPU architecture

To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: **Little Endian**



# Step 1: Determine the CPU architecture

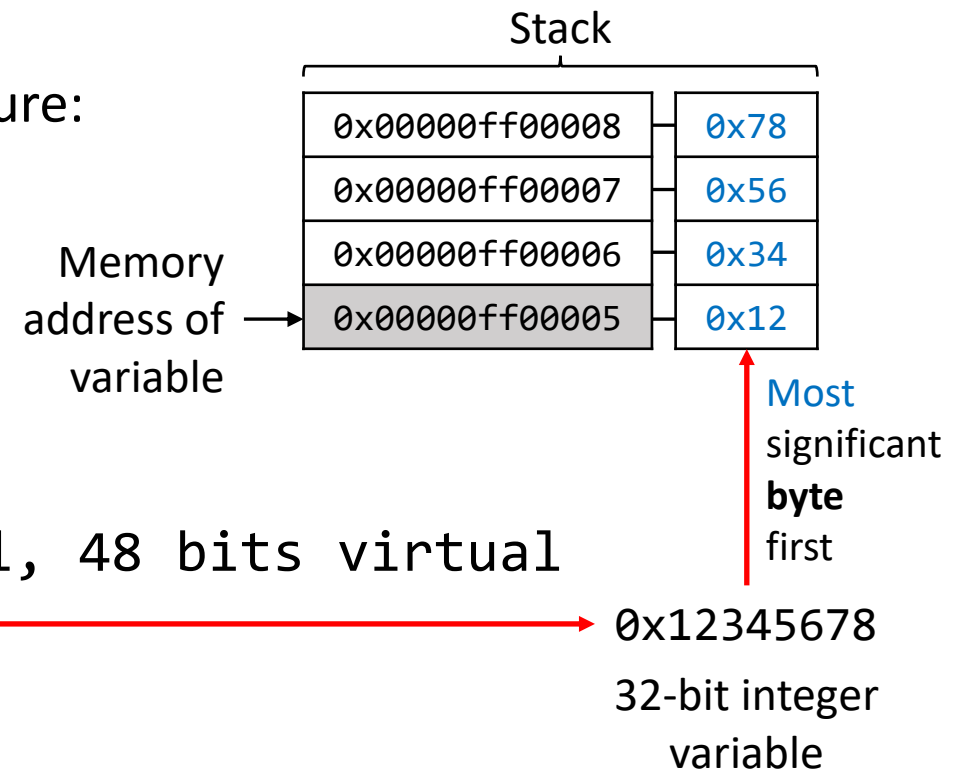
To display information about the CPU architecture:

`lscpu`

The output will look like:

Architecture: x86\_64  
CPU op-mode(s): 32-bit, 64-bit  
Address sizes: 45 bits physical, 48 bits virtual  
Byte Order: **Big Endian**

What if?



## Step 2: Develop code to examine

SCC131\_example2.c

```
1  #include <stdio.h>
2
3  void func1(int a) {
4      int j=6;
5      printf("entered func1");
6  }
7
8  int main(void) {
9      int i=5;
10     func1(7);
11     printf("hello world\n");
12 }
```

- This program will be examined using GDB.
- Notice that the program consists of two functions: `main()` and `func1()`.
- Function `main()` is the caller and `func1()` is the callee.
- Both functions contain local variables.
- Breakpoints will be introduced in **lines 5 and 11**, and the stack frames will be displayed.

```
chatzige@vdi-scclab-001: ~/h-drive/SCC131/code
chatzige@vdi-scclab-001:~/h-drive$ cd SCC131/code
chatzige@vdi-scclab-001:~/h-drive/SCC131/code$ code SCC131_example2.c
chatzige@vdi-scclab-001:~/h-drive/SCC131/code$
```

```
SCC131_example2.c - Visual Studio Code
File Edit Selection View Go Run Terminal Help
Restricted Mode is intended for safe code browsing. Trust this wi... Manage Learn More X

C SCC131_example2.c x
home > chatzige > h-drive > SCC131 > code > C SCC131_example2.c
1  #include <stdio.h>
2
3  void func1(int a) {
4      int j=6;
5      printf("entered func1");
6  }
7
8  int main(void) {
9      int i=5;
10     func1(7);
11     printf("hello world\n");
12 }
```

# SCC Labs



Home



## Step 3: Compile and run code using GDB

---

```
gcc -g SCC131_example2.c -o SCC131_example2
```

← Generate debug information (-g)  
to be used by GDB and create  
executable file (-o).

## Step 3: Compile and run code using GDB

---

```
gcc -g SCC131_example2.c -o SCC131_example2  
gdb SCC131_example2  
(gdb) █
```

Debug the executable file from  
the command line.

## Step 3: Compile and run code using GDB

```
gcc -g SCC131_example2.c -o SCC131_example2
```

```
gdb SCC131_example2
```

```
(gdb) break 5
```

```
Breakpoint 1 at 0x117c: file SCC131_example2.c, line 5.
```

```
(gdb) break 11
```

```
Breakpoint 2 at 0x11b0: file SCC131_example2.c, line 11.
```

```
(gdb) █
```

Introduce  
breakpoints in  
lines 5 and 11.

# Step 3: Compile and run code using GDB

```
gcc -g SCC131_example2.c -o SCC131_example2
```

```
gdb SCC131_example2
```

```
(gdb) break 5
```

```
Breakpoint 1 at 0x117f: file SCC131_example2.c, line 5.
```

```
(gdb) break 11
```

```
Breakpoint 2 at 0x11b3: file SCC131_example2.c, line 11.
```

```
(gdb) run
```

```
Starting program: /home/chatzige/h-drive/SCC131/code/SCC131_example2
```

```
[Thread debugging using libthread_db enabled]
```

```
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
```

```
Breakpoint 1, func1 (a=7) at SCC131_example2.c:5
```

```
5          printf("entered func1");
```

```
(gdb) █
```

Run the executable file until the first breakpoint (in line 5) is reached.

## Step 3: Examine the stack using GDB

---

```
(gdb) info stack  
#0 func1 (a=7) at SCC131_example2.c:5  
#1 0x0000555555551b3 in main () at SCC131_example2.c:10
```


Display the call  
stack (same as  
[backtrace](#))

## Step 3: Examine the stack using GDB

(gdb) `info stack`

#0 `func1` (`a=7`) at `SCC131_example2.c:5`

#1 `0x0000555555551b3` in `main` () at `SCC131_example2.c:10`



```
1  #include <stdio.h>
2
3  void func1(int a) {
4      int j=6;
5      printf("entered func1");
6  }
7
8  int main(void) {
9      int i=5;
10     func1(7);
11     printf("hello world\n");
12 }
```


# Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

In this example:




# Step 3: Examine the stack using GDB

```
(gdb) info stack
```


```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

In this example:



If we had N nested functions:



Frame #0 always denotes the **current** frame.

The indices of frames are updated whenever execution of a function is completed.



## Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 0
```

```
Stack frame at 0x7fffffffdf60:
```

```
rip = 0x55555555517f in func1 (SCC131_example2.c:5); saved rip = 0x5555555551b3
```

```
called by frame at 0x7fffffffdf80
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0x7fffffffdf60
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Display information  
about stack frame 0,  
which corresponds  
to func1.

## Step 3: Examine the stack using GDB

```
(gdb) info stack
#0 func1 (a=7) at SCC131_example2.c:5
#1 0x00005555555551b3 in main () at SCC131_exa
```

```
(gdb) info frame 0
```

**Stack frame at 0x7fffffffdf60:**

```
rip = 0x5555555517f in func1 (SCC131_example2.c:5)
called by frame at 0x7fffffffdf80
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0x7fffffffdf60
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Address that signifies the beginning of this frame (frame 0) and the end of the previous frame (frame 1). When execution of func1 is completed, this frame will be popped from the stack and this address will become the new stack pointer.

## Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 0
```

```
Stack frame at 0x7fffffffdf60:
```

```
rip = 0x55555555517f in func1 (SCC131_example2.c:5); saved rip = 0x5555555551b3
```

```
called by frame at 0x7fffffffdf80
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0x7fffffffdf50
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Address of the *next* instruction pointer, i.e., the address that the frame will return to when operation of func1 resumes (line 5 in func1).

## Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 0
```

```
Stack frame at 0x7fffffffdf60:
```

```
rip = 0x5555555517f in func1 (SCC131_example2.c:5); saved rip = 0x5555555551b3
```

```
called by frame at 0x7fffffffdf80
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Address (and location in the stack) that this frame will return to at the end of func1, when control returns to the previous frame (immediately *after* the call of func1 in line 10).

## Step 3: Examine the stack using GDB

```
(gdb) info stack
#0 func1 (a=7) at SCC131_example2.c:5
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 0
```

```
Stack frame at 0x7fffffffdf60:
```

```
rip = 0x55555555517f in func1 (SCC131_example2.c:5); saved rip = 0x5555555551b3
```

```
called by frame at 0x7fffffffdf80 ←
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0x7fffffffdf60
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Address of the frame  
that called this frame,  
i.e., address of frame 1  
(the caller).

# Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 0
```

```
Stack frame at 0x7fffffffdf60:
```

```
rip = 0x55555555517f in func1 (SCC131_example2.c:5); saved rip = 0x5555555551b3  
called by frame at 0x7fffffffdf80
```

```
source language c.
```

```
Arglist at 0x7fffffffdf50, args: a=7
```

```
Locals at 0x7fffffffdf50, Previous frame's sp is 0x7ff...
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf50, rip at 0x7fffffffdf58
```

Location of the base pointer  
(reference for this frame that  
points to the base pointer of  
the previous frame).

## Step 3: Examine the stack using GDB

```
(gdb) info stack
```

```
#0 func1 (a=7) at SCC131_example2.c:5
```

```
#1 0x00005555555551b3 in main () at SCC131_example2.c:10
```

```
(gdb) info frame 1
```

```
Stack frame at 0x7fffffffdf80:
```

```
rip = 0x5555555551b3 in main (SCC131_example2.c:10); saved rip = 0x7ffff7c29d90
```

```
caller of frame at 0x7fffffffdf60
```

```
source language c.
```

```
Arglist at 0x7fffffffdf70, args:
```

```
Locals at 0x7fffffffdf70, Previous frame's sp is 0x7fffffffdf80
```

```
Saved registers:
```

```
rbp at 0x7fffffffdf70, rip at 0x7fffffffdf78
```

Display information  
about stack frame 1,  
which corresponds  
to main.

## Step 4: Visualise the stack

→ (gdb) `x/28xw $sp`

0x7fffffffdf30:	0x00000000	0x00000000	0x00000000	0x00000007
0x7fffffffdf40:	0x00000000	0x00000000	0x00000000	0x00000006
0x7fffffffdf50:	0xffffdf70	0x0007ffff	0x555551b3	0x00055555
0x7fffffffdf60:	0x00000000	0x00000000	0x00000000	0x00000005
0x7fffffffdf70:	0x00000001	0x00000000	0xf7c29d90	0x0007ffff
0x7fffffffdf80:	0x00000000	0x00000000	0x5555196	0x0005555


8 digits = 32 bits = 4 bytes at 0x7fffffffdf80      4 bytes at 0x7fffffffdf84      4 bytes at 0x7fffffffdf88      4 bytes at 0x7fffffffdf8c

Look at the memory and display (x) the 28 words (w) [note: a *word* contains 32 bits] in hexadecimal format (x) located just below the top of the stack, specified by the stack pointer (\$sp).



# Step 4: Visualise the stack (manually)

```
0x7fffffffdf84: 0x00000000
0x7fffffffdf80: 0x00000000 ----- Beginning of Frame 1 (main) -----
0x7fffffffdf7c: 0x00007fff }
0x7fffffffdf78: 0xf7c29d90 } rip at 0x7fffffffdf78 is 0x00007ffff7c29d90 (saved IP)
0x7fffffffdf74: 0x00000000 }
0x7fffffffdf70: 0x00000001 } rbp at 0x7fffffffdf70 is 0x0000000000000001 (saved BP)
0x7fffffffdf6c: 0x00000005 } i=5 (locally declared variable)
0x7fffffffdf68: 0x00000000
0x7fffffffdf64: 0x00000000
0x7fffffffdf60: 0x00000000 ----- Beginning of Frame 0 (func1) -----
0x7fffffffdf5c: 0x00005555 }
0x7fffffffdf58: 0x555551b3 } rip at 0x7fffffffdf58 is 0x00005555555551b3 (saved IP)
0x7fffffffdf54: 0x00007fff }
0x7fffffffdf50: 0xffffdf70 } rbp at 0x7fffffffdf50 is 0x00007fffffffdf70 (saved BP)
0x7fffffffdf4c: 0x00000006 } j=6 (locally declared variable)
0x7fffffffdf48: 0x00000000
0x7fffffffdf44: 0x00000000
0x7fffffffdf40: 0x00000000
0x7fffffffdf3c: 0x00000007
0x7fffffffdf38: 0x00000000
0x7fffffffdf34: 0x00000000
0x7fffffffdf30: 0x00000000 } sp at 0x7fffffffdf30 (top of stack)
```



SCC131\_example2.c

```
1  #include <stdio.h>
2
3  void func1(int a) {
4      int j=6;
5      printf("entered func1");
6  }
7
8  int main(void) {
9      int i=5;
10     func1(7);
11     printf("hello world\n");
12 }
```

# Step 4: Visualise the stack (manually)

```
0x7fffffffdf84: 0x00000000
0x7fffffffdf80: 0x00000000 ----- Beginning of Frame 1 (main) -----
0x7fffffffdf7c: 0x00007fff }
0x7fffffffdf78: 0xf7c29d90 } rip at 0x7fffffffdf78 is 0x00007ffff7c29d90 (saved IP)
0x7fffffffdf74: 0x00000000 }
0x7fffffffdf70: 0x00000001 } rbp at 0x7fffffffdf70 is 0x0000000000000001 (saved BP)
0x7fffffffdf6c: 0x00000005 } i=5 (locally declared variable)
0x7fffffffdf68: 0x00000000
0x7fffffffdf64: 0x00000000
0x7fffffffdf60: 0x00000000 ----- Beginning of Frame 0 (func1) -----
0x7fffffffdf5c: 0x00005555 }
0x7fffffffdf58: 0x555551b3 } rip at 0x7fffffffdf58 is 0x00005555555551b3 (saved IP)
0x7fffffffdf54: 0x00007fff }
0x7fffffffdf50: 0xffffdf70 } rbp at 0x7fffffffdf50 is 0x00007fffffffdf70 (saved BP)
0x7fffffffdf4c: 0x00000006 } j=6 (locally declared variable)
0x7fffffffdf48: 0x00000000
0x7fffffffdf44: 0x00000000
0x7fffffffdf40: 0x00000000
0x7fffffffdf3c: 0x00000007 } a=7 (function parameter)
0x7fffffffdf38: 0x00000000
0x7fffffffdf34: 0x00000000
0x7fffffffdf30: 0x00000000 } sp at 0x7fffffffdf30 (top of stack)
```

Why is it here? ----- Shouldn't it be here?

# Step 5: Check the registers

---

- In the **x86** architecture, all function parameters appear at the beginning of the relevant stack frame.
- In the **x86-64** architecture, often abbreviated as **x64**, the **first six function parameters** are passed in registers and treated as local variables. The **remaining** function parameters appear at the beginning of the stack frame.
- Type `info reg` to print the content of the registers and notice that value 7 has been stored in register `%rdi`.

Register	Purpose
<code>%rax</code>	temp register; return value
<code>%rbx</code>	callee-saved
<code>%rcx</code>	used to pass 4th argument to functions
<code>%rdx</code>	used to pass 3rd argument to functions
<code>%rsp</code>	stack pointer
<code>%rbp</code>	callee-saved; base pointer
<code>%rsi</code>	used to pass 2nd argument to functions
<code>%rdi</code>	used to pass 1st argument to functions
<code>%r8</code>	used to pass 5th argument to functions
<code>%r9</code>	used to pass 6th argument to functions
<code>%r10-r11</code>	temporary
<code>%r12-r15</code>	callee-saved registers

# Step 6: Disassemble the code (func1)

```
(gdb) disass func1
```

Dump of assembler code for function func1:

```
0x000055555555169 <+0>:    endbr64
0x00005555555516d <+4>:    push    %rbp
0x00005555555516e <+5>:    mov     %rsp,%rbp
0x000055555555171 <+8>:    sub     $0x20,%rsp
0x000055555555175 <+12>:   mov     %edi,-0x14(%rbp)
0x000055555555178 <+15>:   movl    $0x6,-0x4(%rbp)
=> 0x00005555555517f <+22>:   lea     0xe7e(%rip),%rax    # 0x555555556004
0x000055555555186 <+29>:   mov     %rax,%rdi
0x000055555555189 <+32>:   mov     $0x0,%eax
0x00005555555518e <+37>:   call    0x55555555070 <printf@plt>
0x000055555555193 <+42>:   nop
0x000055555555194 <+43>:   leave
0x000055555555195 <+44>:   ret
```

End of assembler dump.

Recall:

```
(gdb) info frame 0
```

Stack frame at 0x7fffffffdf60:

```
rip = 0x5555555517f in func1 (SCC131_example2.c:5);
```

# Step 6: Disassemble the code (main)

(gdb) `disass main`

Dump of assembler code for function main:

```
0x000055555555196 <+0>:  endbr64
0x00005555555519a <+4>:  push   %rbp
0x00005555555519b <+5>:  mov    %rsp,%rbp
0x00005555555519e <+8>:  sub    $0x10,%rsp
0x0000555555551a2 <+12>: movl   $0x5,-0x4(%rbp)
0x0000555555551a9 <+19>: mov    $0x7,%edi
0x0000555555551ae <+24>: call   0x55555555169 <func1>
0x0000555555551b3 <+29>: lea    0xe58(%rip),%rax    # 0x555555556012
0x0000555555551ba <+36>: mov    %rax,%rdi
0x0000555555551bd <+39>: call   0x55555555060 <puts@plt>
0x0000555555551c2 <+44>: mov    $0x0,%eax
0x0000555555551c7 <+49>: leave
0x0000555555551c8 <+50>: ret
```

End of assembler dump.

Recall:

(gdb) `info frame 1`

Stack frame at 0x7fffffffdf80:

`rip = 0x555555551b3` in main (`SCC131_example2.c:10`);

# Step 6: Disassemble the code (main)

(gdb) `disass main`

Dump of assembler code for function main:

```
0x000055555555196 <+0>:  endbr64
0x00005555555519a <+4>:  push   %rbp
0x00005555555519b <+5>:  mov    %rsp,%rbp
0x00005555555519e <+8>:  sub    $0x10,%rsp
0x0000555555551a2 <+12>: movl   $0x5,-0x4(%rbp)
0x0000555555551a9 <+19>: mov    $0x7,%edi
0x0000555555551ae <+24>: call   0x55555555169 <func1>
0x0000555555551b3 <+29>: lea    0xe58(%rip),%rax    # 0x555555556012
0x0000555555551ba <+36>: mov    %rax,%rdi
0x0000555555551bd <+39>: call   0x55555555060 <puts@plt>
0x0000555555551c2 <+44>: mov    $0x0,%eax
0x0000555555551c7 <+49>: leave
0x0000555555551c8 <+50>: ret
```

End of assembler dump.

Evidence that value 7 has been moved to register %rdi – this is the input parameter to func1.

Note that %rdi is a 64-bit register. We use %edi to access the 32 least significant bits of %rdi, given that 7 has been stored as a 32-bit integer.

# Memory layout – Bringing it all together

SCC131\_example3.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int i = 5;
5
6  int func1(int a)
7  {
8      int k = 5;
9  }
10
11 int main()
12 {
13     char * buffer;
14     buffer = (char*)malloc(i+1);
15     if (buffer == NULL) exit(1);
16     func1(i);
17     return 0;
18 }
```

(gdb) `break 9`

Breakpoint 1 at 0x117b: file SCC131\_example3.c, line 9.

(gdb) `run`

Starting program: /home/chatzige/h-drive/SCC131/code/SCC131\_example3

[Thread debugging using libthread\_db enabled]

Using host libthread\_db library "/lib/x86\_64-linux-gnu/libthread\_db.so.1".

Breakpoint 1, func1 (a=5) at SCC131\_example3.c:9

9            }

(gdb) `print &k`

\$1 = (int \*) 0x7fffffffdf3c

(gdb) `print &i`

\$2 = (int \*) 0x555555558010 <i>

(gdb) `frame 1`

#1 0x0000555555551bf in main () at SCC131\_example3.c:16

16            func1(i);

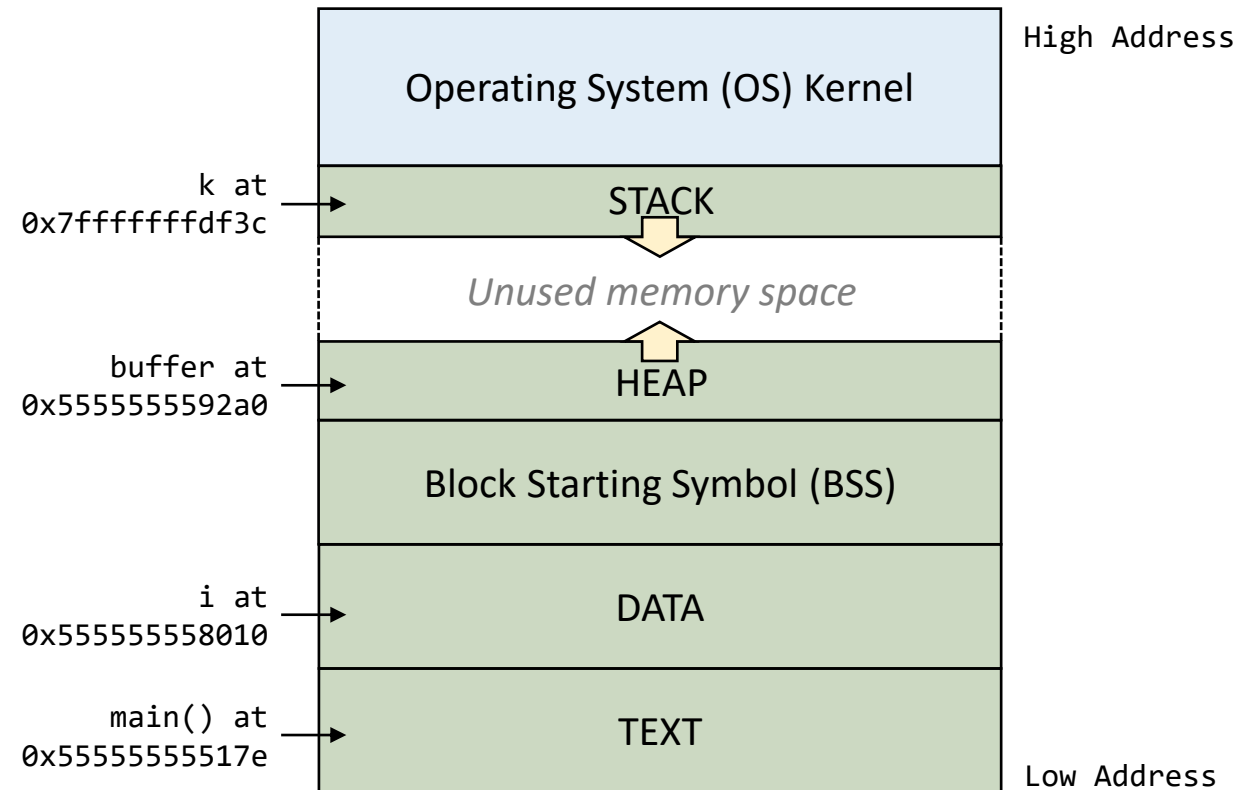
(gdb) `print buffer`

\$3 = 0x5555555592a0 ""

# Memory layout – Bringing it all together

SCC131\_example3.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int i = 5;
5
6  int func1(int a)
7  {
8      int k = 5;
9  }
10
11 int main()
12 {
13     char * buffer;
14     buffer = (char*)malloc(i+1);
15     if (buffer == NULL) exit(1);
16     func1(i);
17     return 0;
18 }
```





# Confused?

---

Unfortunately, memory layout is anything but simple:

- Depends on architecture (Intel, AMD, ARM, ...)
- Depends on OS (macOS, Linux, Windows, ...)
- Depends on compiler and compiler version
- Depends on compiler flags

# Summary

---

We focused on the memory layout of the x86-64 architecture and discussed about:

- The memory management unit, the translation lookaside buffer and the relationship between physical and virtual addresses.
- The main memory regions: Text, Data, BSS, Heap, Stack and OS Kernel.
- The stack layout: the base pointer, the instruction pointer, the stack pointer and the memory regions for function parameters and locally declared variables.
- GDB commands that can help you examine and visualise the stack of a running C program.

# Resources

---

- MIT: [Procedures and stacks](#)
- MIT: [Examining the stack](#)
- MIT: [X86-64 Architecture guide](#)
- Stanford University: [Runtime stack](#) (Lab 7)
- Brown University: [x64 cheat sheet](#) (page 1)
- The Linux Kernel: [Complete virtual memory map with 4-level page tables](#)
- [Stack frame layout on x86-64](#) by Eli Bendersky
- Visual GDB: [GDB command reference](#)